

Lidang Jiang

Birth Date
1995/08/12

Career Focus
MLsys
AI
Embedded Software Engineer
Embodied Intelligence
Robotics
Hardware Engineer

Compensation
30k-35k/mo × 15 monthly salaries

Email
765984881@qq.com (preferred)
2021223075161@alu.scu.edu.cn

Status
Employed; start date negotiable

Website
<https://lidang-jiang.pages.dev>



Location
Xianning, Hubei; based in Shenzhen, Guangdong

Education

Sichuan University, Master's Degree College of Chemical Engineering, Materials and Chemical Engineering - University Second-class Academic Scholarship, Academic Year 2021 - 2022 - University Second-class Academic Scholarship, Academic Year 2022 - 2023 - University Second-class Academic Scholarship, Academic Year 2023 - 2024	2021.09 - 2024.07
Jinzhong University, Bachelor's Degree College of Chemistry and Chemical Engineering, Applied Chemistry	2014.09 - 2018.07
Tongshan County No.1 High School	2011.09 - 2014.07

Self-Evaluation

Scored 502 in College Entrance Exam. Undergraduate GPA: 2.9/4, ranked 45/50. Graduate Entrance Exam: 263 in 2020 and 400 in 2021 (8th/120), Mathematics II: 141 (1st/120). Graduate GPA: 3.5/4, top 30%.

During my graduate studies, I primarily focused on research at the intersection of AI and lithium-ion batteries (LiBs). As the first author, I published two papers in *Energy* and *Applied Energy* (both are Q1 journals in CAS and JCR rankings). Throughout my master's studies, I developed the ability to independently conduct interdisciplinary research.

After completing my master's degree, I have continued to build hands-on experience in AI systems engineering and open-source collaboration. I am passionate about using AI to build everything from inference systems to embodied robots, and my current interests include embodied AI, robotics, inference acceleration, AI infrastructure, embedded software, and hardware engineering.

Publications

- Jiang, Lidang, et al. "A Robust Adapted Flexible Parallel Neural Network Architecture for Early Prediction of Lithium Battery Lifespan." *Energy*, 308:132840, 2024. (IF=10.0)
- Jiang, Lidang, et al. "Generating Comprehensive Lithium Battery Charging Data with Generative AI." *Applied Energy*, 377:124604, 2025. (IF=12.16)

Professional Skills

MLsys / AI Infrastructure Python, PyTorch, SGLang, vLLM / vLLM-Kunlun, CUDA, inference debugging, and environment diagnostics	Languages Python, C++, Go, TypeScript
Infrastructure and Engineering Kubernetes, Docker, Linux, Git	Supplementary Web Engineering Vue, JavaScript, Spring Boot, MySQL, Nginx, Element-UI, Vant

Work Experience

ManpowerGroup Enterprise Management Consulting (Shanghai) Co., Ltd., Software Development Engineer (Embedded Systems) Department: OpenHarmony Enablement Department (Device BG) Base: Bantian, Shenzhen Compensation: CNY 19,500/month 14–16 monthly salaries	2026.07 - Present
Shanghai MicroPort Software Co., Ltd. Shenzhen Branch, Python Backend Developer (AI Inference Framework) Compensation: CNY 16,500/month 12 monthly salaries first 3 months paid at 80%	2025.08 - 2026.07

Hainan Xin Zhujiang Human Resources Development & Management Co., Ltd.	2024.07 - 2025.06
- Large Language Model Algorithm Engineer (2024.07 - 2024.09) - Front-end Development Engineer (2024.09 - 2025.06) Compensation: CNY 8,000/month 12 monthly salaries first 2 months paid at 80%	
Wenjia Education, etc., Graduate Entrance Exam Subject Tutor	2021.05 - 2021.12
Guangzhou Lizhi Education Technology Co., Ltd., Yousi Tutoring, etc., Math Teacher	2021.05 - 2021.12
Others (e.g., Unemployed, Preparing for Exams at Home, etc.)	2019.09 - 2021.05
Real Estate Consultant, Environmental Engineer, Hotel Receptionist, Game Booster	2018.07 - 2019.09

I Project Experience

Project Overview	Project Count
AI Infrastructure, Model Serving, and Training Optimization	15
Embodied AI, Robotics, and Reinforcement Learning	7
Systems Engineering, Developer Tooling, and Industrial Integration	5
Web, Government, and Industry Digitalization	10
Total	37

I AI Infrastructure, Model Serving, and Training Optimization

Project 1: vLLM-Kunlun (Baidu Kunlun Chip Inference Framework)	2025.08.18 —2026.07.06
Type: Open-Source Contribution + Product Delivery Role: Contributor / MLsys Engineer Repo: baidu/vLLM-Kunlun GitHub Stars: 461. vLLM inference framework for Kunlun XPU; before open-sourcing I helped narrow a Qwen2.5-7B Attention failure on P800 Float16, and after open-sourcing I landed 9 merged PRs: fixed an InternVL inference KeyError (PR #108), improved installation and tutorial docs (PR #136), added <code>collect_env.py</code> diagnostics (PR #218), and covered 0.15.1 release alignment, Read the Docs cleanup, Docker documentation maintenance, and XProfiler benchmark documentation fixes (PR #249, PR #251, PR #253, PR #327, PR #328, PR #341).	
Project 2: FunASR	2026.05.19
Type: Open-Source Contribution Role: Contributor Repo: modelscope/FunASR GitHub Stars: 19,728. End-to-end speech recognition toolkit and open-source pretrained model library; fixed an <code>UnboundLocalError</code> when VAD is enabled without <code>punc_model</code> by initializing <code>punc_res</code>, and added <code>punc_model=None</code> regression tests (PR #2840).	
Project 3: vLLM	2026.05.05
Type: Open-Source Contribution Role: Contributor Repo: vllm-project/vllm GitHub Stars: 88,561. High-throughput LLM inference and serving engine; improved tokenizer-class error handling by recommending Transformers upgrades when model tokenizers require newer library support (PR #38099).	
Project 4: OpenClaw	2026.04.28 —2026.07.29
Type: Open-Source Contribution Role: Contributor Repo: openclaw/openclaw GitHub Stars: 385,590. Cross-platform personal AI assistant; landed 2 merged PRs: added explicit chat type support for active-memory configuration and runtime gating with schema and execution regression coverage (PR #66285), and made the Gmail watcher exclude SPAM, TRASH, DRAFT, and SENT labels to prevent sent mail and drafts from entering inbound hook processing, with regression coverage (PR #56720).	
Project 5: ExecuTorch	2026.04.10 —2026.05.08
Type: Open-Source Contribution Role: Contributor Repo: pytorch/executorch GitHub Stars: 4,876. On-device AI runtime for PyTorch; landed 3 merged PRs adding dynamic shape tests for xnnpack models, Linux Python runtime end-to-end XNNPACK tests, and a <code>torch.split</code> export fix for alias annotations (PR #18701, PR #18703, PR #18700).	
Project 6: LitServe	2026.04.07 —2026.04.09
Type: Open-Source Contribution Role: Contributor Repo: Lightning-AI/LitServe GitHub Stars: 3,923. AI inference server framework; landed 2 merged PRs fixing an auth bypass in multi-API deployments plus health-check exception handling, stale state, and concurrency behavior (PR #674, PR #673).	
Project 7: Olive	2026.04.08
Type: Open-Source Contribution Role: Contributor Repo: microsoft/Olive GitHub Stars: 2,377. Model finetuning, conversion, and quantization toolkit; fixed external data loading in HQQ and RTN quantization passes (PR #2380).	
Project 8: anomalib	2026.04.08
Type: Open-Source Contribution Role: Contributor Repo: open-edge-platform/anomalib GitHub Stars: 6,025. Anomaly detection library; landed 3 merged PRs fixing <code>pandas>=3.0 DataFrame</code> enum compatibility, PatchCore coreset initial random point handling, and deprecated API / resource leak issues (PR #3512, PR #3509, PR #3508).	
Project 9: PyTorch AO	2026.04.07
Type: Open-Source Contribution Role: Contributor Repo: pytorch/ao GitHub Stars: 2,932. PyTorch quantization and optimization toolkit; skipped <code>linear+bn</code> fusion for inputs above rank 2 in PT2E flows to prevent invalid graph rewrites (PR #4242).	
Project 10: SGLang	2026.04.06
Type: Open-Source Contribution Role: Contributor Repo: sgl-project/sglang GitHub Stars: 31,555. High-performance LLM and multimodal serving framework; landed 2 merged PRs adding auth/token unit tests and regression coverage for Hermes, Llama32, and Mistral <code>function_call</code> detectors (PR #21400, PR #21399).	

Project 11: RLLM

2026.04.04

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [rllm-org/rllm](#) | **GitHub Stars:** 5,771. RL-for-LLM research framework; updated ver1 0.7.1+ import path compatibility after upstream path changes (PR #480).

Project 12: vLLM-Omni

2026.03.05 —2026.05.06

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [vllm-project/vllm-omni](#) | **GitHub Stars:** 5,976. Omni/multimodal inference extension on top of vLLM; landed 4 merged PRs covering error-response HTTP status handling, missing ANSI colors in piped CLI logo output, optional gradio packaging, and replacing bare `print()` calls with structured logging (PR #1687, PR #1636, PR #2221, PR #2228).

Project 13: OpenRLHF

2026.04.03

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [OpenRLHF/OpenRLHF](#) | **GitHub Stars:** 9,897. Distributed RLHF training framework; ensured Ray runtime respects user-defined `NCCL_DEBUG` values instead of overwriting them (PR #1212).

Project 14: Baidu AIK-SGLang DeepSeek / GLM5 PD Deployment

2026.03.17 —2026.07.06

- **Type:** Closed-Source Baidu AI Inference System | **Role:** MLsys / Inference Engineer. Built and maintained Kubernetes FedDeployment-based Prefill/Decode-disaggregated deployments for DeepSeek-V4-Flash and GLM5 on Kunlun P800/P900 clusters with AIK-customized SGLang. Split prefill, decode, and router launch flows, tuning multi-node TP/DP/EP, DeepEP, chunked prefill, DP attention, and cache-report settings. Added XPU profiler / CPU-affinity runtime patching, per-rank profiler-device injection, `bench_serving --pd-separated` profiling, and `aiakperf` long-context benchmark scripts to make Prefill/Decode performance debugging and capacity checks reproducible.

Project 15: Transformers

2026.03.30

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [huggingface/transformers](#) | **GitHub Stars:** 163,484. Hugging Face's flagship model library; removed an incorrect `torchvision` dependency from PIL backend image processors to reduce unnecessary installs and runtime failures (PR #45045).

I Embodied AI, Robotics, and Reinforcement Learning

Project 1: rclcpp

2026.06.08

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [ros2/rclcpp](#) | **GitHub Stars:** 787. ROS 2 C++ client library; backported the already-merged rolling/kilted fix to Jazzy by initializing `GenericClient` result response offsets, reducing uninitialized-offset build and runtime risk (PR #3139).

Project 2: Launch ROS

2026.05.10

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [ros2/launch_ros](#) | **GitHub Stars:** 82. ROS 2 launch tooling for starting and testing ROS nodes; rejected the deprecated node-name frontend key to keep launch descriptions aligned with current ROS 2 node naming behavior (PR #538).

Project 3: Gymnasium

2026.04.12

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [Farama-Foundation/Gymnasium](#) | **GitHub Stars:** 12,301. Standard RL environment API; added a `RepeatAction` wrapper for repeated-action execution (PR #1553).

Project 4: HighwayEnv

2026.04.07 —2026.06.20

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [Farama-Foundation/HighwayEnv](#) | **GitHub Stars:** 3,298. Autonomous-driving RL environment library; landed 2 merged PRs fixing `SubprocVecEnv ConnectionResetError` failures under `forkserver/spawn` start methods and `neighbour_vehicles` misses on connected lanes (PR #668, PR #667).

Project 5: Genesis

2026.03.28 —2026.04.05

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [Genesis-Embodied-AI/genesis-world](#) | **GitHub Stars:** 29,718. General-purpose embodied AI simulation platform; landed 6 merged PRs across short-path motion planning, multi-robot IK DOF mapping, camera lookat behavior, rigid-body energy APIs, WSL2 GPU detection, and propeller API naming cleanup (PR #2609, PR #2610, PR #2612, PR #2613, PR #2614, PR #2653).

Project 6: PyTorch RL

2026.04.05

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [pytorch/rl](#) | **GitHub Stars:** 3,512. Reinforcement learning toolkit in the PyTorch ecosystem; added a `strict_shape` parameter to `QValueModule` for stronger action-shape enforcement (PR #3593).

Project 7: ManiSkill

2026.03.30

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [mani-skill/ManiSkill](#) | **GitHub Stars:** 3,209. Robotics manipulation simulator and benchmark suite; landed 2 merged PRs fixing `replay_trajectory` compatibility with `gymnasium>=1.0` and a `PlugCharger-v1` dense-reward doc label (PR #1402, PR #1403).

I Systems Engineering, Developer Tooling, and Industrial Integration

Project 1: Kubernetes

2026.06.15 —2026.07.01

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [kubernetes/kubernetes](#) | **GitHub Stars:** 124,377. Production-grade container orchestration system; landed 3 merged PRs migrating `Node.Spec.ProviderID` set-once constraints, `StatefulSet` immutable-field validation, and `HPA CrossVersionObjectReference.Kind/Name` validation to the declarative validation path, with equivalence coverage and generated validation updates across affected API versions (PR #138080, PR #138077, PR #138082).

Project 2: crawllee-python

2026.04.14

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [apify/crawllee-python](#) | **GitHub Stars:** 9,414. Python crawling and browser-automation library; replaced fixed sleeps with poll-based assertions in a flaky autoscaling test to improve CI stability (PR #1835).

Project 3: opcua-asyncio	2026.04.05
• Type: Open-Source Contribution Role: Contributor Repo: FreeOpcUa/opcua-asyncio GitHub Stars: 1,453. Asyncio-based OPC UA implementation; landed 4 merged PRs fixing explicit certificate-format handling, uint16 sequence-number overflow, Nodetyped typing, and PubSub TTL / key-name validation (PR #1954, PR #1955, PR #1956, PR #1957).	
Project 4: Docker Compose	2026.04.03
• Type: Open-Source Contribution Role: Contributor Repo: docker/compose GitHub Stars: 38,035. Multi-service container orchestration tool; returned non-ErrNotExist stat failures from Tar.Sync() instead of swallowing them silently (PR #13684).	
Project 5: Everything Claude Code	2026.03.28 —2026.03.29
• Type: Open-Source Contribution Role: Contributor Repo: affaan-m/ECC GitHub Stars: 238,856. Agent configuration and hooks automation toolkit; landed 5 merged PRs covering Windows path compatibility, Codex persistent_instructions, MCP startup_timeout_sec, observer temp-file handling, and hooks denylist migration (PR #971, PR #972, PR #974, PR #977, PR #992).	
I Web, Government, and Industry Digitalization	
Project 1: University Researcher Profile System	2025.05.25 —2025.06.04
• Type: Personal Open-Source Full-Stack Web App Role: Independent Full-Stack Developer Repo: Lidang-Jiang/UniversityResearcherProfiles GitHub Stars: 0. Personal portal system for research groups; independently built both front-end clients and backend REST APIs, added bilingual switching, and deployed it to a live lab website.	
Project 2: Haikou 12345 Intelligent Report System V1.0	2025.05.13 —2025.05.31
• Type: Government Hotline Automation System Role: Front-End Developer. Built the intelligent template management flow, interactive report editor, and DeepSeek prompt dialog for structured report generation from work-order data.	
Project 3: Government-Enterprise Platform (V2 / V3.1)	2025.02.21 —2025.05.31
• Type: Government Enterprise Data Platform Role: Front-End Developer. Delivered Excel batch import, enterprise verification, data-status tracking, and batch approval capabilities for cross-department enterprise-data workflows.	
Project 4: 928 Hospital Model Command System V2.0	2025.03.31 —2025.04.09
• Type: Military Hospital Information System Role: Front-End Developer. Built the to-do and task-supervision module with dual task views and a department tree selector.	
Project 5: Risk Monitoring and Early Warning System Common Service Platform V1.0.1	2025.03.10 —2025.03.21
• Type: Government Risk Warning Platform Role: Front-End Developer. Built the warning-model management module with three-state filtering, model/entity linkage, and automatic warning-name population.	
Project 6: Red Detachment of Women Memorial Park Digital Platform (V1.0 / V1.2.0)	2024.11.11 —2025.03.06
• Type: Cultural Tourism Digital Platform Role: Front-End Developer. Iterated the admin panel, youth-facing client, and reservation workflow, covering study-tour registration, visit services, secure document downloads, and server-side Word template rendering.	
Project 7: Project Platform Management System V2.6	2025.01.13 —2025.02.06
• Type: Enterprise Internal Management System Role: Mobile Front-End Developer. Built the mobile approval center for business-opportunity, overtime, contract, and related workflows.	
Project 8: Digital Campus System V2.0.5	2024.12.13 —2024.12.24
• Type: Education Information Platform Role: Front-End Developer. Built the PC-side work-assignment module plus WeChat-side delegation verification and push workflows with QR-code verification.	
Project 9: Project Operation Management Platform V2.5	2024.10.22 —2024.11.20
• Type: Enterprise Internal Management System Role: Front-End Developer. Delivered PC-side vehicle reports and mobile workbench, approval-center, and vehicle-management modules.	
Project 10: Haikou Police Administrative Knowledge Base	2024.07.25 —2024.09.25
• Type: Government AI Application Role: Algorithm Engineer. Participated in administrative data cleaning, RAG-system construction, and LLM evaluation, analyzing responses to 20,000+ real government queries to improve answer quality.	