

Lidang Jiang

Birth Date
1995/08/12

Phone
15307241568 (limited phone availability during work hours)

Email
765984881@qq.com (preferred)
2021223075161@alu.scu.edu.cn

Location
Xianning, Hubei; based in Shenzhen, Guangdong

Title
MLsys - Embodied Intelligence - AI - Robotics - Embedded Software Engineer

Compensation
25k-30k/mo

Website
<https://lidang-jiang.pages.dev>

Availability
2 weeks



Education

Sichuan University, Master's Degree College of Chemical Engineering, Materials and Chemical Engineering - University Second-class Academic Scholarship, Academic Year 2021 - 2022 - University Second-class Academic Scholarship, Academic Year 2022 - 2023 - University Second-class Academic Scholarship, Academic Year 2023 - 2024	2021.09 - 2024.07
Jinzhong University, Bachelor's Degree College of Chemistry and Chemical Engineering, Applied Chemistry	2014.09 - 2018.07
Tongshan County No.1 High School	2011.09 - 2014.07

Self-Evaluation

Scored 502 in College Entrance Exam. Undergraduate GPA: 2.9/4, ranked 45/50. Graduate Entrance Exam: 263 in 2020 and 400 in 2021 (8th/120), Mathematics II: 141 (1st/120). Graduate GPA: 3.5/4, top 30%.

During my graduate studies, I primarily focused on research at the intersection of AI and lithium-ion batteries (LiBs). As the first author, I published two papers in *Energy* and *Applied Energy* (both are Q1 journals in CAS and JCR rankings). Throughout my master's studies, I developed the ability to independently conduct interdisciplinary research.

After completing my master's degree, I have continued to build hands-on experience in AI systems engineering and open-source collaboration. I am passionate about using AI to build everything from inference systems to embodied robots, and my current interests include embodied AI, robotics, inference acceleration, AI infrastructure, and Embedded Software.

Publications

- 1:** Jiang, Lidang, et al. "A Robust Adapted Flexible Parallel Neural Network Architecture for Early Prediction of Lithium Battery Lifespan." *Energy*, 308:132840, 2024. (IF=10.0)
- 2:** Jiang, Lidang, et al. "Generating Comprehensive Lithium Battery Charging Data with Generative AI." *Applied Energy*, 377:124604, 2025. (IF=12.16)

Professional Skills

MLsys / AI Infrastructure
Python, PyTorch, SGLang, vLLM / vLLM-Kunlun, CUDA, inference debugging, and environment diagnostics

Infrastructure and Engineering
Kubernetes, Docker, Linux, Git

Languages
Python, C++, Go, TypeScript

Supplementary Web Engineering
Vue, JavaScript, Spring Boot, MySQL, Nginx, Element-UI, Vant

Work Experience

Shanghai MicroPort Software Co., Ltd. Shenzhen Branch, Python Backend Developer	2025.08 - Present
Hainan Xin Zhujiang Human Resources Development & Management Co., Ltd. Large Language Model Algorithm Engineer (2024.07 - 2024.09)	2024.07 - 2025.06

• Front-end Development Engineer (2024.09 - 2025.06)

Wenjia Education, etc., Graduate Entrance Exam Subject Tutor	2021.05 - 2021.12
Guangzhou Lizhi Education Technology Co., Ltd., Yousi Tutoring, etc., Math Teacher	2021.05 - 2021.12
Others (e.g., Unemployed, Preparing for Exams at Home, etc.)	2019.09 - 2021.05
Real Estate Consultant, Environmental Engineer, Hotel Receptionist, Game Booster	2018.07 - 2019.09

Project Experience

AI Infrastructure, Model Serving, and Training Optimization

| Project 1: vLLM-Kunlun (Baidu Kunlun Chip Inference Framework) 2025.08.18 —Present

- **Type:** Open-Source Contribution + Product Delivery | **Role:** Contributor / MLsys Engineer | **Repo:** [baidu/vLLM-Kunlun](#) | **GitHub Stars:** 442. vLLM inference framework for Kunlun XPU; before open-sourcing I helped narrow a Qwen2.5-7B Attention failure on P800 Float16, and after open-sourcing I landed 9 merged PRs: fixed an InternVL inference KeyError ([PR #108](#)), improved installation and tutorial docs ([PR #136](#)), added `collect_env.py` diagnostics ([PR #218](#)), and covered 0.15.1 release alignment, Read the Docs cleanup, Docker documentation maintenance, and XProfiler benchmark documentation fixes ([PR #249](#), [PR #251](#), [PR #253](#), [PR #327](#), [PR #328](#), [PR #341](#)).

| Project 2: FunASR 2026.05.19

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [modelscope/FunASR](#) | **GitHub Stars:** 18,756. End-to-end speech recognition toolkit and open-source pretrained model library; fixed an `UnboundLocalError` when VAD is enabled without `punc_model` by initializing `punc_res`, and added `punc_model=None` regression tests ([PR #2840](#)).

| Project 3: vLLM 2026.05.05

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [vllm-project/vllm](#) | **GitHub Stars:** 84,966. High-throughput LLM inference and serving engine; improved tokenizer-class error handling by recommending Transformers upgrades when model tokenizers require newer library support ([PR #38099](#)).

| Project 4: OpenClaw 2026.04.28

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [openclaw/openclaw](#) | **GitHub Stars:** 381,242. Cross-platform personal AI assistant; added explicit chat type support for active-memory configuration and runtime gating with schema and execution regression coverage ([PR #66285](#)).

| Project 5: ExecuTorch 2026.04.10 —2026.05.08

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [pytorch/executorch](#) | **GitHub Stars:** 4,771. On-device AI runtime for PyTorch; landed 3 merged PRs adding dynamic shape tests for xnnpack models, Linux Python runtime end-to-end XNNPACK tests, and a `torch.split` export fix for alias annotations ([PR #18701](#), [PR #18703](#), [PR #18700](#)).

| Project 6: LitServe 2026.04.07 —2026.04.09

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [Lightning-AI/LitServe](#) | **GitHub Stars:** 3,900. AI inference server framework; landed 2 merged PRs fixing an auth bypass in multi-API deployments plus health-check exception handling, stale state, and concurrency behavior ([PR #674](#), [PR #673](#)).

| Project 7: Olive 2026.04.08

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [microsoft/Olive](#) | **GitHub Stars:** 2,343. Model finetuning, conversion, and quantization toolkit; fixed external data loading in HQQ and RTN quantization passes ([PR #2380](#)).

| Project 8: anomalib 2026.04.08

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [open-edge-platform/anomalib](#) | **GitHub Stars:** 5,892. Anomaly detection library; landed 3 merged PRs fixing pandas>=3.0 DataFrame enum compatibility, PatchCore coreset initial random point handling, and deprecated API / resource leak issues ([PR #3512](#), [PR #3509](#), [PR #3508](#)).

| Project 9: PyTorch AO 2026.04.07

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [pytorch/ao](#) | **GitHub Stars:** 2,881. PyTorch quantization and optimization toolkit; skipped `linear+bn` fusion for inputs above rank 2 in PT2E flows to prevent invalid graph rewrites ([PR #4242](#)).

| Project 10: SGLang 2026.04.06

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [sgl-project/sglang](#) | **GitHub Stars:** 29,856. High-performance LLM and multimodal serving framework; landed 2 merged PRs adding auth/token unit tests and regression coverage for Hermes, Llama32, and Mistral `function_call` detectors ([PR #21400](#), [PR #21399](#)).

| Project 11: RLLM 2026.04.04

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [rllm-org/rllm](#) | **GitHub Stars:** 5,666. RL-for-LLM research framework; updated `verl 0.7.1+` import path compatibility after upstream path changes ([PR #480](#)).

| Project 12: vLLM-Omni 2026.03.05 —2026.05.06

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [vllm-project/vllm-omni](#) | **GitHub Stars:** 5,382. Omni/multimodal inference extension on top of vLLM; landed 4 merged PRs covering error-response HTTP status handling, missing ANSI colors in piped CLI logo output, optional gradio packaging, and replacing bare `print()` calls with structured logging ([PR #1687](#), [PR #1636](#), [PR #2221](#), [PR #2228](#)).

| Project 13: OpenRLHF 2026.04.03

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [OpenRLHF/OpenRLHF](#) | **GitHub Stars:** 9,723. Distributed RLHF training framework; ensured Ray runtime respects user-defined `NCCL_DEBUG` values instead of overwriting them ([PR #1212](#)).

| Project 14: Baidu AIAK-SGLang DeepSeek / GLM5 PD Deployment 2026.03.17 —2026.06.24

- **Type:** Closed-Source Baidu AI Inference System | **Role:** MLsys / Inference Engineer. Built and maintained Kubernetes FedDeployment-based Prefill/Decode-disaggregated deployments for DeepSeek-V4-Flash and GLM5 on Kunlun P800/P900 clusters with AIAK-customized SGLang. Split prefill, decode, and router launch flows, tuning multi-node TP/DP/EP, DeepEP, chunked prefill, DP attention, and cache-report settings. Added XPU profiler / CPU-affinity runtime patching, per-rank profiler-device injection, `bench_serving --pd-separated` profiling, and `aiakperf` long-context benchmark scripts to make Prefill/Decode performance debugging and capacity checks reproducible.

| Project 15: Transformers 2026.03.30

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [huggingface/transformers](#) | **GitHub Stars:** 162,064. Hugging Face's flagship model library; removed an incorrect `torchvision` dependency from PIL backend image processors to reduce unnecessary installs and runtime failures (PR #45045).

Embodied AI, Robotics, and Reinforcement Learning

| Project 1: rclcpp 2026.06.08

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [ros2/rclcpp](#) | **GitHub Stars:** 770. ROS 2 C++ client library; backported the already-merged rolling/kilted fix to Jazzy by initializing `GenericClient` result response offsets, reducing uninitialized-offset build and runtime risk (PR #3139).

| Project 2: Launch ROS 2026.05.10

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [ros2/launch_ros](#) | **GitHub Stars:** 80. ROS 2 launch tooling for starting and testing ROS nodes; rejected the deprecated node-name frontend key to keep launch descriptions aligned with current ROS 2 node naming behavior (PR #538).

| Project 3: Gymnasium 2026.04.12

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [Farama-Foundation/Gymnasium](#) | **GitHub Stars:** 12,107. Standard RL environment API; added a `RepeatAction` wrapper for repeated-action execution (PR #1553).

| Project 4: HighwayEnv 2026.04.07 —2026.06.20

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [Farama-Foundation/HighwayEnv](#) | **GitHub Stars:** 3,283. Autonomous-driving RL environment library; landed 2 merged PRs fixing `SubprocVecEnv ConnectionResetError` failures under `forkserver/spawn` start methods and `neighbour_vehicles` misses on connected lanes (PR #668, PR #667).

| Project 5: Genesis 2026.03.28 —2026.04.05

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [Genesis-Embodied-AI/genesis-world](#) | **GitHub Stars:** 29,466. General-purpose embodied AI simulation platform; landed 6 merged PRs across short-path motion planning, multi-robot IK DOF mapping, camera lookat behavior, rigid-body energy APIs, WSL2 GPU detection, and propeller API naming cleanup (PR #2609, PR #2610, PR #2612, PR #2613, PR #2614, PR #2653).

| Project 6: PyTorch RL 2026.04.05

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [pytorch/rl](#) | **GitHub Stars:** 3,476. Reinforcement learning toolkit in the PyTorch ecosystem; added a `strict_shape` parameter to `QValueModule` for stronger action-shape enforcement (PR #3593).

| Project 7: ManiSkill 2026.03.30

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [mani-skill/ManiSkill](#) | **GitHub Stars:** 3,056. Robotics manipulation simulator and benchmark suite; landed 2 merged PRs fixing `replay_trajectory` compatibility with `gymnasium>=1.0` and a `PlugCharger-v1` dense-reward doc label (PR #1402, PR #1403).

Systems Engineering, Developer Tooling, and Industrial Integration

| Project 1: Kubernetes 2026.06.15 —2026.06.26

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [kubernetes/kubernetes](#) | **GitHub Stars:** 123,298. Production-grade container orchestration system; landed 2 merged PRs migrating `StatefulSet` immutable-field validation and `HPA CrossVersionObjectReference.Kind/Name` validation to the declarative validation path, with equivalence coverage and generated validation updates across affected API versions (PR #138080, PR #138077).

| Project 2: crawlee-python 2026.04.14

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [apify/crawlee-python](#) | **GitHub Stars:** 9,272. Python crawling and browser-automation library; replaced fixed sleeps with poll-based assertions in a flaky autoscaling test to improve CI stability (PR #1835).

| Project 3: opcua-asyncio 2026.04.05

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [FreeOpcUa/opcua-asyncio](#) | **GitHub Stars:** 1,439. Asyncio-based OPC UA implementation; landed 4 merged PRs fixing explicit certificate-format handling, `uint16` sequence-number overflow, `NodeId` typing, and `PubSub` TTL / key-name validation (PR #1954, PR #1955, PR #1956, PR #1957).

| Project 4: Docker Compose 2026.04.03

- **Type:** Open-Source Contribution | **Role:** Contributor | **Repo:** [docker/compose](#) | **GitHub Stars:** 37,691. Multi-service container orchestration tool; returned `non-ErrNotExist` stat failures from `Tar.Sync()` instead of swallowing them silently (PR #13684).

